
Formatting Instructions For NeurIPS 2026

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The abstract paragraph

2 **1 Introduction**

3 **2 Methods and Setup**

4 **2.1 Direct Preference Optimization (DPO)**

- 5 • **Dataset:** UltraFeedback Preference Pairs (course provided)
- 6 • **Model:** Qwen2.5-1.5B-Instruct (Forzen with LoRA)

7 **3 Results**

8 **3.1 Direct Preference Optimization (DPO)**

9 **3.1.1 Standard DPO**

10 **3.2 Proximal Policy Optimization (PPO)**

11 **3.3 Group Relative Policy Optimization (GRPO)**

12 **3.4 Safety Calibration of the Aligned Policies**

13 **3.5 Feedback Source: RLVR versus RLAIIF**

14 **3.6 Cross Task Discussion**

15 **4 Conclusions**

16 **5 References**

17 **Appendix**

18 **A More Method Details**

19 **B Deviations and Caveats**

20 **B.1 Direct Preference Optimization (DPO)**

- 21 • We changed the Batch Size to 1 and Gradient Accumulate Steps to 16, keeping effective
- 22 batch size constant. This was done as our hardware would run out of memory otherwise.